ASSIGNMENT

1:table created:

create table temperature_data

(dat String,

zip INT,

temp INT)

row format delimited fields terminated by ',';

data loaded:

load data local inpath 'location' into table temperature_data;

```
hive> select *  from temperature_data;
OK
10-01-1990      123112  10
14-02-1991      283901  11
10-03-1990      381920  15
10-01-1991      302918  22
12-02-1990      384902  9
10-01-1991      123112  11
14-02-1990      283901  12
10-03-1991      381920  16
10-01-1990      302918  23
12-02-1991      384902  10
10-01-1993      123112  11
14-02-1994      283901  12
10-03-1993      381920  16
10-01-1994      302918  23
12-02-1991      384902  10
10-01-1991      123112  11
14-02-1990      283901  12
10-03-1991      381920  16
10-01-1990      302918  23
12-02-1991      384902  10
```

2:zip code greater than 300000 and less than 399999

```
hive> select dat,temp from temperature_data where zip>300000
    > ;
OK
10-03-1990      15
10-01-1991      22
12-02-1990      9
10-03-1991      16
10-01-1990      23
12-02-1991      10
10-03-1993      16
10-01-1994      23
12-02-1991      10
10-03-1991      16
10-01-1990      23
12-02-1991      10
Time taken: 2.861 seconds, Fetched: 12 row(s)
```

3:max temperature:

```
    > select MAX(temp) from temperature_data;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider u
tion engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20190113182322_f964cbc6-0f42-4e15-b962-b8b0bbf50680
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1547372468711_0005, Tracking URL = http://localhost:8088/proxy/application_154737246
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job  -kill job_1547372468711_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-01-13 18:24:25,644 Stage-1 map = 0%,  reduce = 0%
2019-01-13 18:25:01,930 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 4.79 sec
2019-01-13 18:25:30,660 Stage-1 map = 100%,  reduce = 67%, Cumulative CPU 8.75 sec
2019-01-13 18:25:33,946 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 10.97 sec
MapReduce Total cumulative CPU time: 10 seconds 970 msec
Ended Job = job_1547372468711_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 10.97 sec   HDFS Read: 8325 HDFS Write: 102 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 970 msec
OK
23
```

4:export view to file in local file system

Command : insert overwrite local directory ‘location’ row format delimited fields terminated by ‘| ‘

```
hive> insert overwrite local directory '/home/acadgild/Hbase' row format delimited fields terminated by '|'
    > select * from temperature_data_vw;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execu
tion engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20190113183314_255224b6-96c3-4327-89ae-fdc6d5077038
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1547372468711_0006, Tracking URL = http://localhost:8088/proxy/application_1547372468711_0006/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job  -kill job_1547372468711_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-01-13 18:33:53,074 Stage-1 map = 0%,  reduce = 0%
2019-01-13 18:34:20,192 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 5.23 sec
2019-01-13 18:34:52,580 Stage-1 map = 100%,  reduce = 67%, Cumulative CPU 9.64 sec
2019-01-13 18:34:55,567 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 11.78 sec
MapReduce Total cumulative CPU time: 11 seconds 780 msec
Ended Job = job_1547372468711_0006
Moving data to local directory /home/acadgild/Hbase
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 11.78 sec   HDFS Read: 7976 HDFS Write: 3 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 780 msec
OK
Time taken: 104.614 seconds
```